AI AGENT GOVERNANCE

5 Signs Your AI Agents Aren't Audit-Ready

The compliance gaps most teams don't see until an auditor asks.

AI Identity

ai-identity.co

SIGN 1 OF 5

1

You're Logging Events, Not Evidence

Application logs ≠ audit trail. If your agent records are rows in Splunk or CloudWatch, they can be altered without detection. Auditors need tamper-evident records with chain-of-custody guarantees — not searchable log files.

SIGN 2 OF 5

2

Your Agents Share API Keys

Shared credentials make it impossible to attribute a specific action to a specific agent. Non-repudiation requires unique, cryptographically-bound identity per agent — not a platform-managed shared secret.

SIGN 3 OF 5

3

You Can't Reconstruct a Decision

An auditor asks: 'Why did Agent X take Action Y at 2:47 AM?' If you can't replay the complete decision path — policy evaluation, tool calls, downstream responses — you have observability, not forensics.

SIGN 4 OF 5

4

Compliance Is a Manual Spreadsheet

If proving EU AI Act or SOC 2 compliance means pulling screenshots and filling Excel templates, you're doing it at O(n) cost per audit. Automated compliance assessments with exportable evidence packages reduce this to O(1).

SIGN 5 OF 5

5

Auditors Have to Trust Your Infrastructure

If your evidence can only be verified inside your platform, auditors are trusting the vendor — not the evidence. Independent verification means any third party can validate your audit chain without accessing your systems.

THE SOLUTION

The Fix

01 **Register**

Issue unique identity per agent

02 **Enforce**

Deny-by-default policy gateway

03 **Prove**

HMAC-SHA256 hash-chained forensic evidence

Register → Enforce → Prove

Start Your Free Trial

ai-identity.co

Unique agent identity · Policy enforcement · Forensic evidence

Get Started Free

AI Identity

AI Agent Governance Platform