

EXECUTIVE R&D REPORT

Agent Accountability Standards

Landscape, AI Identity's position, and the next bet

The identity and authorization layer for AI agents is consolidating fast — and every standard in it defers the tamper-evident, non-repudiable audit layer. That omission is the bet.

PREPARED FOR

Jeff Leva, Founder — AI Identity

GENERATED BY

standards-architect engine (Claude Code, Opus 4.8)

DATE

2026-06-09 · grounded in live GitHub + web research

The bet: become the non-repudiation layer of record

Standardization of AI agents has split into two layers. The identity + authorization layer — who an agent is and what it may do — is consolidating quickly across MCP, A2A, the OAuth-for-agents stack, and OpenID's accountability group. The forensic layer — tamper-evident, signed, replayable proof of what an agent actually did — is named as a requirement by regulators and governance bodies, but no standard carries it. That gap is durable, dated, and unclaimed.

RECOMMENDATION

Make AI Identity the non-repudiation layer of record for agent activity: drive the OCSF generic non-repudiation / tamper-evident-logging profile (AI as first consumer) and ship the conformant emitter — timed to the EU AI Act record-keeping clauses that bind from Aug 2026. The precondition just landed: our serialization primitive (#1662) merged today.

IMMEDIATE NEXT MOVE (NOT THE DESTINATION)

Get #1661 (attestation) over its review bar, and post a SHORT supportive note on Ania's freshly-filed #1665 (delegation) proposing the attestation-to-delegation seam as a question to her — listen-first. File the profile only after #1661 proves the pattern and the room signals. The profile is the prize; the seam is the step.

Why now — three things converged this week:

- Our signature-canonicalization primitive (#1662) MERGED into the OCSF development line — the first piece of the forensics agenda to actually land.
- Ania filed the delegation object (#1665) — the seam we'd been holding for is now live and proposable, with her.
- The EU AI Act's logging / record-keeping obligations for high-risk systems bind from Aug 2026, with the technical standards still in draft — leadership room.

State of play

OCSF — the schema layer AI Identity profiles as its forensic envelope:

- AI observability shipped in OCSF 1.8.0 (Mar 2026): `ai_operation` profile, `ai_model`, `message_context`. The foundation is in a tagged release.
- Agent-identity layer is in OPEN PRs: #1641 (`ai_agent`) and #1665 (delegation), both Ania's (Cisco). Umbrella issue #1640 defines the full design — a new AI category (`uid 9`, not `7`), control-plane classes (`agent_activity`, `delegation_activity`), and a delegation-lineage graph — none filed yet.
- Non-repudiation wedge (ours): #1662 (serialization on `digital_signature`) MERGED 2026-06-09 to the dev line the maintainer calls '1.9.0-dev' — no 1.9.0 tag exists yet. #1661 (attestation object) is OPEN, needs its approvals.
- The maintainer roadmap (AWS) is silent on attestation / non-repudiation / forensics — this is a bottom-up wedge, which is both the opening and the risk.

The wider field — identity consolidating, audit deferred:

- MCP and A2A were both donated to the Linux Foundation's new Agentic AI Foundation (Dec 2025 / Jun 2025). Platinum members: AWS, Anthropic, Block, Bloomberg, Cloudflare, Google, Microsoft, OpenAI.
- Auth is solved-ish: MCP on OAuth 2.1 + Resource Indicators (no audit section); A2A via Agent Cards (TLS + bearer tokens, no signed action log); the OAuth-for-agents stack (ID-JAG / Cross-App Access, Transaction Tokens, on-behalf-of) races to model delegated authority chains.
- Governance venues are explicitly asking for accountability: NIST's NCCoE concept paper names 'logging and transparency' and 'access delegation for accountability'; OpenID's AIIM group lists accountability in its top-three scope.

THE WHITE SPACE, CONFIRMED BY OMISSION

No identity or authorization standard carries normative tamper-evident or non-repudiation requirements. OpenTelemetry GenAI — the de-facto agent telemetry convention — explicitly disclaims the immutable / audit-ready layer. Nobody has filled this; the venues are asking who will.

SOURCES

- 1 OCSF: gh PR status #1641/#1661/#1662/#1665 (2026-06-09); [ocsf/ocsf-schema](#) issue #1640; AWS Open Source blog, OCSF 1.8.0 / ITU support (Mar 2026).
- 2 Linux Foundation press, Agentic AI Foundation + A2A project (Dec 2025 / Jun 2025); [modelcontextprotocol.io](#) authorization spec 2025-11-25; [a2a-protocol.org](#) v0.3.0.
- 3 NIST CAISI AI Agent Standards Initiative + NCCoE concept paper (Feb 2026); OpenID AIIM Community Group; [opentelemetry.io](#) GenAI semantic conventions.

Where we stand

- Two things our product already emits map cleanly onto the schema: the gateway `audit_log` (immutable per-action rows, HMAC-SHA256 hash chain, per-org chains) maps to the attestation object; the Mandate Service (signed delegation grants, ECDSA-P256 over RFC 8785 / JCS, hybrid classical+PQC signature array) maps to delegation + `digital_signature`.
- We hold the canonical-serialization primitive (#1662, merged) that the whole non-repudiation chain signs over — and the attestation object (#1661) that binds agent identity to a signed, tamper-evident record. Both are deliberately domain-agnostic: AI is the first consumer, not the only one.
- Framing authority: AI Identity anchored the 'AI forensics as its own concern' boundary in the working group, and was offered OCSF reviewer (approver) status — shaping power over what AI-related schema lands, not just contributor status.

THE MOAT IS ORDERING, NOT FEATURES

Non-repudiation grows upward from an identity root and a canonical serialization — it cannot be retrofitted onto telemetry after the fact. Detection/response tools (e.g. AgentDR) are downstream consumers of a non-repudiable record, not competitors to it. Whoever owns the signing-from-identity layer is upstream of the entire accountability stack.

Two honest caveats, attached:

- Concentration / framing risk: the agentic-AI schema effort currently rests on two contributors (Ania + Jeff) with no visible big-vendor co-sponsor on these specific PRs. 'Define the standard' is the upside — IF a co-sponsor and the review process hold; the PRs could still be reshaped.
- 'Merged' discipline: only #1662 is merged (to the dev line; no 1.9.0 release tag). Externally, say 'contributing to OCSF' and cite #1662 as the merged proof point — do not claim the rest has landed.

SOURCES

- 4 AI Identity repo: `gateway audit_log` (`common/audit`), Mandate Service (`mandate/app`, ECDSA-P256 / JCS, hybrid PQC array); `memory project_ocsf_engagement`, `project_mandate_service`.
- 5 OCSF PR #1662 (merged 2026-06-09), PR #1661 (open); working-group provenance notes (2026-05-29 → 06-09).

The three open lanes

Ranked by what we can act on and where being wrong is costliest — not by novelty.

- 1

Non-repudiation profile + conformant emitter (OCSF)

ACT NOW

The category-defining move. Precondition #1662 merged; #1665 (delegation) just filed, so the attestation-to-delegation seam is now proposable. NIST/OpenID/EU all asking for accountability. Lane: ours (#1661) + a seam with Ania (#1665) — propose with her, never redefine her object.
- 2

Post-quantum-signed evidence records

BUILD-IN

Differentiator and standards-forward. ML-DSA (FIPS 204) is shipping in CAs and crypto libraries, yet no agent standard mandates or even profiles PQ signatures for credentials or audit records — and audit records need decades-long integrity (harvest-now, verify-later). Aligns with the H2 Mandate Service hybrid design already on the books. Lower urgency, ours to claim.
- 3

Regulatory mapping (EU AI Act / NIST NCCoE)

FRAME

Highest demand-generation, dated to Q4 2026. Map the forensic-audit capability to the AI Act's record-keeping / logging / traceability clauses and NCCoE's accountability pillar. This is the wedge into regulated verticals — the medical full-audit-chain use case a third party raised in the WG is the lead example. More GTM/compliance than R&D, but it is what turns the standard into revenue.

THESE ARE COMPLEMENTARY, NOT EXCLUSIVE

Lane 1 wins the seat and ships the reference implementation; Lane 3 supplies the dated demand that justifies it; Lane 2 is the differentiator you bake into Lane 1's records rather than a separate project. Sequence, don't choose.

The bet, checked against our stack

Make AI Identity the non-repudiation layer of record for agent activity — drive the OCSF generic non-repudiation profile and ship the conformant emitter, timed to the EU AI Act Q4 2026 forcing function.

- ✓ **Operational surface — low**
Additive to services we already run (gateway, Mandate). No new always-on infrastructure, which matters on GCP-credit runway.
- ✓ **Dogfoodable**
We already emit the primitives on our own dashboard's verified agents — the emitter proves itself against our own audit trail.
- ✓ **Solo-runnable**
Standards work + an emitter profile is one-person-shippable; it does not require a team to operate, unlike a detection/response product.
- ✓ **Crypto-agility**
The signature array (classical + reserved ML-DSA slot) is already designed — Lane 2 folds in without rework.
- ~ **Spec alignment — verify**
The profile must match the published spec's L2 session-attestation (DSSE + ES256). Align the profile to the spec, or update the spec with a section number — don't let them drift.

RISKS & CONDITIONS

(1) No big-vendor co-sponsor yet — recruit one before over-investing. (2) EU AI Act technical standards are still in draft — map to them as leadership, but never claim conformance to a standard that isn't published. (3) Lane discipline: the delegation seam is Ania's — the win is a joint proposal, not a unilateral one.

VERDICT

Build. The bet survives every stack-fit check, the precondition merged this week, the venue is asking, and the regulatory clock is running. The only gating move is social, not technical: clear #1661, and open the seam with Ania on #1665.