

AI IDENTITY — Every agent leaves a trace.

FOR TEAMS THAT RUN AGENTS, NOT TRAIN MODELS

The EU AI Act Deployer Evidence Checklist

The records Articles 12 and 26 require, the questions a market surveillance authority will ask, and a worksheet to find out where each record lives today.

Why now

On 2 August 2026 the high-risk obligations under Annex III and the deployer duties under Article 26 became applicable. The six-month log-retention clock is already running — records cannot be backfilled.

August 2026 edition · Regulation (EU) 2024/1689

Engineering aid, not legal advice. Verify article citations against the OJ text with counsel.

You run agents. That makes you a deployer — and the deployer duties get the least airtime.

Most EU AI Act content is written for model providers: GPAI documentation, training data summaries, conformity assessments. If your company runs AI agents rather than trains models, you are a deployer under the Act. Your obligations are shorter to list — and they are the ones your team will be asked to evidence first.

The staged timeline

- 2 Feb 2025 Prohibited practices apply.
- 2 Aug 2025 General-purpose AI model rules apply.
- 2 Aug 2026 **General application: Annex III high-risk obligations and the Article 26 deployer duties. This is the one that lands on agent operators.**
- 2 Aug 2027 Annex I product-safety high-risk systems.

What Article 26 actually asks of you

- Art. 26(1)** Operate each system per the provider's instructions — which means a record of how each agent was configured, and when that changed.
- Art. 26(2)** Assign human oversight: a named function with the competence and authority to intervene, and a record that intervention happened when it should have.
- Art. 26(5)** Monitor operation, and inform the provider or authorities when something looks like a serious incident.
- Art. 26(6)** Retain the automatically generated logs for at least six months — longer where other EU or national law says so.

Behind these sits Article 12: high-risk systems must automatically record events over their lifetime, at a level that makes the system's behavior traceable. That is where the records in this checklist come from.

Why ordinary logs don't clear the bar

A market surveillance authority — or your own counsel, after an incident — asks about a decision one of your agents made in March:

"How do you know this record wasn't edited after the fact?"

If the answer involves the phrase “our vendor's dashboard,” that isn't evidence — it's a relationship. Application logs in a SaaS console are mutable rows in someone else's database. They can be complete, well-intentioned, and beautifully searchable, and still prove nothing, because nothing about them lets a third party verify that the sequence of events is intact. If you can't verify it offline, it's not evidence.

Five questions to ask before an authority does

- 1 Can you produce March's records without logging in to any vendor's dashboard?
- 2 Can a third party verify the sequence of events is intact — offline, without trusting you or your vendor?
- 3 Which of your agents fall under an Annex III high-risk category? Is that a query today, or a quarterly archaeology project?
- 4 Who approved the last high-impact agent action — and where is the approval, denial, or expiry recorded?
- 5 If your logging vendor disappeared tomorrow, would the evidence you owe under Article 26(6) survive?

A “no” on any of these is not a compliance-tool gap. It is a records gap — and records, unlike policy documents, cannot be written late.

Ten records the Act expects you to produce

For each record: write down where it lives today, then answer the only question that makes it evidence — could someone verify it without trusting the system that produced it?

☐

Automatic event records

Art. 12(1)

Every agent decision recorded automatically over the system's lifetime — not just errors.

Where it lives today: _____ Verifiable offline? Y / N

☐

Traceable agent identity

Art. 12

Each record names which agent acted, for whom, and with what outcome.

Where it lives today: _____ Verifiable offline? Y / N

☐

Tamper-evidence

Art. 12(4)

A third party can verify offline that the record sequence is intact (e.g. hash-chained events, signed attestations).

Where it lives today: _____ Verifiable offline? Y / N

☐

Configuration record

Art. 26(1)

How each agent was configured against the provider's instructions — and when that changed.

Where it lives today: _____ Verifiable offline? Y / N

☐

Human oversight record

Art. 14 / 26(2)

A named oversight function, plus the full approval trail: approved, denied, and auto-expired.

Where it lives today: _____ Verifiable offline? Y / N

Ten records the Act expects you to produce (cont.)



Monitoring & incidents

Art. 26(5)

Operational monitoring with an escalation path to the provider or authorities for serious incidents.

Where it lives today: _____ Verifiable offline? Y / N



Log retention

Art. 19 / 26(6)

Automatically generated logs kept at least six months; longer where other law applies.

Where it lives today: _____ Verifiable offline? Y / N



Risk classification inventory

Annex III

An Annex III category (or explicit not-in-scope) attached to every agent, queryable on demand.

Where it lives today: _____ Verifiable offline? Y / N



Capability disclosures

Art. 13

What each agent could do at each point in the period — descriptions and capability sets over time.

Where it lives today: _____ Verifiable offline? Y / N



Policy change history

Art. 9

The rules agents ran under, with their full change history as risk management evolves.

Where it lives today: _____ Verifiable offline? Y / N

WHAT GOOD LOOKS LIKE

One bundle, one article per artifact, verifiable without us

This is the artifact set our EU AI Act export profile produces for a reporting period — generated from hash-chained audit records and checkable with an open CLI, no account on our side required:

```
eu_ai_act_2024_export/  
|- manifest.json           DSSE-signed index of the bundle  
|- annex_iv_documentation.json  Annex IV technical documentation  
|- access_log.csv          Art. 12 - every agent decision in period  
|- attestations/<session>.dsse.json Art. 12(4) - signed session attestations  
|- human_oversight_log.csv    Art. 14 / 26(2) - approvals, denials, expiries  
|- agent_risk_classification.csv Annex III category per agent  
|- policy_change_log.csv     Art. 9 - risk-management change history  
|- capability_disclosures.csv Art. 13 - per-agent capabilities over time
```

Two things we deliberately don't claim

Article 10 data governance is the provider's problem.

Upstream training-data governance sits outside a deployer platform's visibility. Our Annex IV export says so explicitly rather than fabricating evidence to fill the section.

There is no EU AI Act certificate to wave.

Conformity assessment bodies are still standing up. What can exist today is architecture that produces the required records in the required shape — and an export a regulator can verify without trusting the vendor.

The six-month clock is already running

The logs you must produce next February are the ones your agents are generating right now. If the worksheet left blanks in the “verifiable offline” column, that is the gap to close first. See how AI Identity produces this evidence set from per-agent identity, fail-closed policy, and hash-chained audit trails:

ai-identity.co · ai-identity.co/eu-ai-act-checklist