

POSITIONING · AGENT-SECURITY POSTURE

Where AI Identity Sits in Anthropic's Agent-Security Model

An honest posture map across the zero-trust tiers — and the layer above them that is ours alone.

Prepared for · AI Identity leadership / partner conversations

Date · June 23, 2026

Basis · Insight #101 · briefing #215 · Sprint 16 item #416

The bar is rising on credentials. We sit above it.

Anthropic is raising the bar on agent credentials. AI Identity already operates at the layer above — the portable record of authority and action — and has now closed the one credential gap that mattered.

The model (our synthesis — not an Anthropic-branded framework)

Anthropic publishes zero-trust *principles*, not a numbered tier ladder. We synthesize them into a 3-tier maturity model to map ourselves honestly. Their published controls (Claude Code security docs): “multiple short-lived, narrowly scoped credentials ... expiring independently”; **fail-closed** defaults; least-privilege permissioning.

Tier	What it requires	Source basis
Tier 1 Identity hygiene	Unique cryptographic identity per agent; short-lived, narrowly-scoped credentials; deny-by-default. <i>Static API keys unsuitable even here.</i>	Claude Code security docs
Tier 2 Transport & key custody	Mutual TLS; keys held in hardware (HSM/TPM); credentials issued via federation (WIF/OIDC), not stored secrets.	Anthropic Workload Identity Federation
Tier 3 Hardware attestation	Remote attestation of the agent runtime (TPM/enclave/confidential compute); signing identity bound to attested hardware.	Anthropic confidential inference (TPM root)

AI Identity posture (honest current state)

Tier	AI Identity status	Evidence
Tier 1	Largely met — just advanced.	Per-agent <code>aid_sk_</code> identity; runtime keys now expire by default + refresh flow (PR #349, in review); gateway fail-closed .
Tier 2	Partial — our root is Tier-2-grade.	Signing root is Cloud KMS / HSM-backed (ECDSA P-256, JWKS published). Federated issuance for customer agent creds is roadmap (#414); mTLS not yet.
Tier 3	Roadmap.	Consume + record TPM/enclave/mTLS attestation at registration, bound into the OCSF evidence chain (#415).

What the tiers don't cover is exactly our product.

Above the tiers — the part Anthropic doesn't occupy

Tiers 1–3 secure *how an agent authenticates*. AI Identity adds the layer above: a **portable, cross-vendor, offline-verifiable record of what authority an agent held and what it did** — OCSF-native audit events plus DSSE/ECDSA-signed evidence a third party can verify with **only our published public key** (no shared secret). Hardware attestation makes this layer *more* valuable: recorded properly, we become the record layer **on top of** hardware roots — feeding the OCSF workload-attestation gap we are co-authoring with IBM.

“Are you obsolete?” — No. Hardware-bound credentials are a **complementary lower layer, not a competitor**. They raise demand for a portable record that can *incorporate* attestations — precisely our layer. The one genuine catch-up — the static agent key Anthropic calls unsuitable at Tier 1 — is **closed** (PR #349). Our signing root is already HSM-backed. We are not behind the trend; we sit on top of it.

Roadmap (Sprint 16 board items)

#413	Short-lived runtime tokens by default — DONE , PR #349 (in review).
#414	OIDC / Workload Identity Federation for customer agent credentials (Tier 2).
#415	Consume + record hardware attestation (Tier 3) — the moat / OCSF workload-attestation gap.
#416	This posture map (positioning).

What to watch for (reviewer's note)

- The **“3 tiers” are our synthesis**, not an Anthropic-labeled framework. Principles are sourced; verify against Anthropic's primary docs before external/customer use.
- Tier 1 **“largely met” is contingent** on PR #349 merging and deploying. Until then say “in review,” not “shipped.”
- Tier 2 **“HSM-backed” refers to our signing/attestation root**, not customer agent-credential custody (still classical, hashed-at-rest).
- **mTLS-for-agent-creds and WIF are roadmap, not shipped** — don't imply otherwise.
- The **“Anthropic doesn't occupy the record layer” claim is our assessment** as of 2026-06; revisit if they ship a provenance/attestation-record product.

Sources

Anthropic — Claude Code Security (short-lived/scoped/expiring credentials; fail-closed; least privilege): code.claude.com/docs/en/security · Anthropic Workload Identity Federation for the Claude API · Anthropic confidential inference / TPM root of trust · Internal: Insight #101, briefing #215, PR #349.