

AI Identity

Cryptographic identity, policy, and audit for AI agents.

The problem

Enterprises are moving fast into agentic AI. Agents are about to take real actions on real systems — create records, move data, call APIs, spend money on the company's behalf.

But every agent today is running on a shared service account key. There's no way to answer three questions that every security, compliance, and platform leader needs answered:

- **Which agent** took this action?
- **Under whose authority** was it authorized?
- **Against which policy** was it evaluated?

The identity, policy, and audit primitives we built for humans and services over the last 20 years simply don't exist for agents. AI Identity closes that gap.

What AI Identity does

- **Cryptographic agent identity.** Every agent gets a signed, revocable, rotatable identity — not a shared key.
 - **Context-aware policy enforcement.** Every agent decision is evaluated against the rules you set, at decision time.
 - **Immutable, auditor-ready audit log.** Each decision is chained and signed using the same primitive git uses for commit history — tamper-evident by construction.
-

Four things a security leader should know

1. Not in the hot path. Agents verify attestations offline. AI Identity signs once; your runtime verifies locally. Zero added latency on agent decisions. Scale your agents to a million calls/sec — we don't care.

2. Hash-chained at the source, normalized at render — your audit log is immutable, your UI is consistent. Every decision is cryptographically linked to the one before it at write time, then rendered through a canonical view layer. The chain holds even when the dashboard layout changes. An auditor can replay the chain a year later and prove it hasn't been tampered with. SOC 2, EU AI Act, and NIST AI RMF export profiles built in.

3. Least-privilege by default, proven. Per-agent, per-decision policy — not per-service-account. The audit trail is the proof you enforced it.

4. Crypto-agile by design. ECDSA-P256 today, ML-DSA-87 slot ready. Your mandates won't need a re-signing event when NIST clocks tick over — the signature envelope already carries an algorithm identifier and a second-signature slot, so hybrid post-quantum signing rolls in without invalidating anything we issued today.

Platform-enforced hardening

You configure once; we enforce every call.

The infrastructure controls below are not customer responsibilities to wire up. They are platform guarantees that ship by default to every tenant.

- **Cluster.** Binary Authorization in ENFORCE mode (only attested, digest-pinned images run). Pod Security Admission at restricted. Default-deny NetworkPolicies with explicit allow-lists. Non-root containers with read-only root filesystems.
- **Ingress.** Cloud Armor WAF in ENFORCE with adaptive DDoS protection and per-IP throttling. HTTPS everywhere. No public kubectl endpoint — control plane reached only through Connect Gateway with Master Authorized Networks on.
- **Secrets.** Google Secret Manager via CSI driver — no secrets in manifests. Signing keys held in Cloud KMS HSM, never leave the HSM boundary. Upstream LLM credentials encrypted with Fernet at rest.
- **Database.** PostgreSQL with Row-Level Security in FORCE mode and sslmode=verify-full. Application-layer scoping is layered on top of database-layer RLS, so a compromised service cannot read another tenant's data.
- **Application.** ABAC on agent metadata. Offline attestation verification CLI. Real-time SIEM push via signed webhook. Compliance export API mapped to SOC 2 CC6/CC7, ISO 27001 A.12/A.13, and CIS Kubernetes Benchmark.

Fit for a security team

Capability	Detail
Agent identity	Issuance, rotation, revocation, TTL
Audit log	Immutable, append-only, cryptographically chained
Policy	Context-aware, evaluated at decision time
Compliance exports	SOC 2, EU AI Act, NIST AI RMF profiles
SIEM integration	Audit export API
Deployment	Your cloud or ours

The design partner ask

- **One agent workflow** you pick. Starts small, expands if it works.
- **30 days, free.** No contract games.
- **Weekly 20-min sync** directly with me.
- **Success criteria** we define together in week 1.
- **In exchange:** honest feedback, and — if it works for you — a reference conversation with one other buyer.

That's it. If it isn't a fit at day 30, we part as friends and I keep buying you coffee.

About the founder

Jeff Leva — Founder & CEO, AI Identity.

I spent years inside the kind of production systems that can't fail — cloud banking platforms handling \$50B+ in client assets, 99.9% uptime, zero tolerance for ambiguity about who did what and why.

When I started working with agentic AI systems, I kept running into the same problem: agents executing real actions with shared keys, no identity, and no audit trail. The governance primitives we'd built for humans — identity, policy, traceability — simply didn't exist for agents.

AI Identity is my answer to that gap: cryptographic agent identity, immutable audit logs, and policy enforcement built to the standards production systems actually require.

Contact

Jeff Leva — Founder & CEO

jeff@ai-identity.co

ai-identity.co